

Sai Narayana B

📞 +91 8248881924

✉️ sainarayana2311@gmail.com

🐙 GitHub: SaiNarayana-B

🌐 LinkedIn: [sai23](#)

🌐 Portfolio: [sai-narayana-portfolio.vercel.app](#)

SUMMARY

GenAI Engineer with 1.5 years of experience building and shipping production-grade AI agent systems at Quantiphi. Reduced multi-agent pipeline runtime by 95% via dynamic parallel orchestration; shipped an open-source codebase analysis tool (patchwork-conventions) to PyPI and MCP Registry. Skilled in multi-agent orchestration, Retrieval Augmented Generation (RAG), NL2SQL, LLM evaluation, and agent SDK integration (Google ADK, Claude, Copilot) on Google Cloud Platform.

EXPERIENCE

Quantiphi

Machine Learning Engineer (Applied GenAI)

Jul 2025 - Present

Bengaluru, India

AI Agent Platform — Enterprise Orchestration (Google ADK & MCP)

- Architected an internal AI agent platform using Google ADK and MCP servers to standardize tool integration, data access, and orchestration for enterprise AI automation systems.
- Implemented dynamic LLM routing and token tracing to select appropriate models by task complexity, improving latency and reducing inference cost.

TEX Agent — Multi-Agent Staffing System (RAG & Parallel Orchestration)

- Developed a multi-agent staffing recommendation system using RAG pipelines over 2 years of historical SOW data with fan-out/fan-in orchestration for any number of roles simultaneously.
- Impact: Reduced pipeline runtime by 95% — from 45 minutes to 5-7 minutes via dynamic parallel orchestration.

NL2SQL Agent — QAssure Platform (Text-to-SQL & MCP)

- Developed a conversational NL2SQL agent translating natural language queries into BigQuery SQL via a BigQuery MCP server, with schema-aware prompting and output summarization.

LLM Caching & Cost Optimization (Prompt Engineering)

- Ran implicit vs. explicit caching POC; restructured all agent and orchestrator system prompts — reducing inference cost by over 80% across production agent runs.

SoW Summarization Pipeline & LLM Wiki (RAG & Knowledge Management)

- Built end-to-end pipeline covering Drive ingestion, LLM summarization, vector indexing, and GCS storage; built REST APIs for LLM Wiki sync to keep the internal knowledge base current.

Keybank SDD Demo (Spec-Driven Development & Microservices)

- Delivered client-facing demo using Spec-Driven Development and harness engineering across a multi-microservices architecture with frontend UI; presented to client stakeholders.

Codeaira — Internal AI Coding Assistant (CLI & VS Code Extension)

- Overhauled CLI and VS Code extension: optimized system prompts, added new tools, invokable slash commands for skills, and extended discoverability to .github, .claude, .agents folders.

Cypher — AI Developer Analytics (VS Code Extension)

- Built a developer tool that captures AI interaction logs from GitHub Copilot, Kiro, and Claude Code; uses LLM-based analysis to evaluate prompt quality and calculate effort-saved karma scores.

Quantiphi

Machine Learning Engineer Intern

Jan 2025 - Jun 2025

Bengaluru, India

AI Canvas Agent & Rapid Pitch Agent (Agentic AI & Google ADK)

- Built the AI Canvas Agent, enabling GTM teams to generate workshop-ready solution documentation by mapping stakeholders, business impact, and ROI using prompt-driven workflows.
- Developed the Rapid Pitch Agent, an AI agent that researches target companies and analyzes financials, AI/ML initiatives, and competitive landscape to generate strategic pitch insights.
- Implemented Google ADK framework and developed MCP servers and tools for integrating enterprise data sources with AI agents across 3 internal pilot projects.

SKILLS

Languages: Python, SQL

AI Systems: Multi-Agent Orchestration, Large Language Models (LLMs), RAG Pipelines, NL2SQL, Prompt Engineering, Prompt Caching, Chain-of-Thought, Few-Shot Prompting, Token Optimization, LLM Evaluation, Agentic AI, MCP Servers, A2A Protocol, Function Calling, Model Routing, Guardrails

Frameworks: Google ADK, LangChain, LangGraph, Claude Agent SDK, Copilot SDK, OpenAI API, Gemini API, Anthropic API, FastAPI, Streamlit

Cloud & Infra: Google Cloud Platform (Vertex AI, BigQuery, GCS, Cloud Run), Docker, REST APIs, GitLab

Tools: VS Code Extension Development, Vector Databases, Embeddings, Git

OPEN SOURCE PROJECT

patchwork-conventions

GitHub · PyPI · MCP Registry · Jun 2026

AST-based codebase scanner that auto-generates CONVENTIONS.md for AI coding agents

- Built tree-sitter AST analysis engine detecting naming conventions, import patterns, error handling, testing frameworks, and API shapes across Python, TypeScript, JavaScript, Go, and Rust with confidence scores and real code examples.

- Shipped as CLI tool, MCP server (8 tools), and Python API; published to PyPI (patchwork-conventions) and MCP Registry; integrates with Claude Code, Cursor, and GitHub Copilot.

EDUCATION

SRM Institute of Science and Technology	2021 - 2025
<i>B.Tech in Computer Science Engineering</i>	<i>CGPA: 9.17</i>

CERTIFICATIONS

- Google Cloud Professional Machine Learning Engineer
- Google Cloud Associate Cloud Engineer
- Google Cloud Generative AI Leader
- Machine Learning Specialization — Stanford University & DeepLearning.AI