

Sai Narayana B

+91 8248881924

sainarayana2311@gmail.com

GitHub: SaiNarayana-B

LinkedIn: sai23

Portfolio: sai-narayana-portfolio.vercel.app

SUMMARY

GenAI Engineer with 1.5 years of experience building and shipping production-grade AI agent systems at Quantiphi. Reduced multi-agent pipeline runtime by 95% via dynamic parallel orchestration; shipped an open-source codebase analysis tool (patchwork-conventions) to PyPI and MCP Registry. Skilled in multi-agent orchestration, Retrieval Augmented Generation (RAG), NL2SQL, LLM evaluation, and agent SDK integration (Google ADK, Claude, Copilot) on Google Cloud Platform.

EXPERIENCE

Quantiphi

Jul 2025 - Present

Machine Learning Engineer (Applied GenAI)

Bengaluru, India

TEX Agent — Multi-Agent Staffing System (RAG & Parallel Orchestration)

- Built a multi-agent staffing system using Google ADK that takes a project ID or role description, dynamically spawns parallel agents for each required role, and returns top candidate recommendations with fit rationale.
- Built a SoW ingestion pipeline that pulls documents from Google Drive, summarizes them via LLM, and indexes into a vector store for the staffing agent's RAG retrieval.
- Reduced pipeline runtime by 95% (45 min → 5 min) by replacing sequential role processing with dynamic parallel orchestration.

NL2SQL Agent — QAssure Platform (Text-to-SQL & MCP)

- Built a conversational agent that lets PMs and business leads query project health data in plain English, converts it to BigQuery SQL via an MCP server, and returns summarized results.

LLM Caching & Cost Optimization (Prompt Engineering)

- Ran implicit vs. explicit caching POC and restructured all agent system prompts, cutting daily inference cost from \$4,000 to \$2,500 (37% reduction).
- Set up dynamic LLM routing to pick cheaper/faster models based on task complexity, reducing inference cost across production agents.

LLM Wiki — Internal Knowledge Base (REST APIs & GCS)

- Built REST APIs to sync the internal knowledge base (LLM Wiki) with GCS, keeping the documents used by production AI agents up to date.

Keybank SDD Demo (Spec-Driven Development & Microservices)

- Delivered a client demo where AI agents generated specs, built a multi-service backend with frontend, tested, and deployed it end-to-end using spec-driven development; presented to KeyBank stakeholders.

Codeaira — Internal AI Coding Assistant (CLI & VS Code Extension)

- Audited and improved the internal AI coding assistant (CLI + VS Code extension) by benchmarking against open-source tools, rewriting system prompts, and adding slash-command based skill invocation.

Cypher — AI Developer Analytics

- Built a developer tool that captures AI interaction logs from Copilot, Kiro, and Claude Code, scores prompt quality using LLM analysis, and calculates effort-saved metrics per developer.

Quantiphi

Jan 2025 - Jun 2025

Machine Learning Engineer Intern

Bengaluru, India

AI Canvas Agent & Rapid Pitch Agent (Agentic AI & Google ADK)

- Built the AI Canvas Agent that generates solution docs and slide decks for GTM teams by mapping stakeholders, business impact, and ROI from a single project brief.
- Built the Rapid Pitch Agent that researches a target company's financials, AI/ML initiatives, and competitors, then produces a ready-to-use pitch deck for pre-sales teams.
- Implemented Google ADK framework and built MCP servers to connect enterprise data sources with AI agents across 3 internal pilot projects.

SKILLS

Languages: Python, SQL

AI Systems: Multi-Agent Orchestration, Large Language Models (LLMs), RAG Pipelines, NL2SQL, Prompt Engineering, Prompt Caching, Chain-of-Thought, Few-Shot Prompting, Token Optimization, LLM Evaluation, Agentic AI, MCP Servers, A2A Protocol, Function Calling, Model Routing, Guardrails

Frameworks: Google ADK, LangChain, LangGraph, Claude Agent SDK, Copilot SDK, OpenAI API, Gemini API, Anthropic API, FastAPI, Streamlit

Cloud & Infra: Google Cloud Platform (Vertex AI, BigQuery, GCS, Cloud Run), Docker, REST APIs, GitLab

Tools: VS Code Extension Development, Vector Databases, Embeddings, Git

OPEN SOURCE PROJECT

patchwork-conventions

GitHub · PyPI · MCP Registry · Jun 2026

AST-based codebase scanner that auto-generates CONVENTIONS.md for AI coding agents

- Built tree-sitter AST analysis engine detecting naming conventions, import patterns, error handling, testing frameworks, and API shapes across Python, TypeScript, JavaScript, Go, and Rust with confidence scores and real code examples.
- Shipped as CLI tool, MCP server (8 tools), and Python API; published to PyPI (patchwork-conventions) and MCP Registry; integrates with Claude Code, Cursor, and GitHub Copilot.

EDUCATION

SRM Institute of Science and Technology

2021 - 2025

B.Tech in Computer Science Engineering

CGPA: 9.17

CERTIFICATIONS

- Google Cloud Professional Machine Learning Engineer
- Google Cloud Associate Cloud Engineer
- Google Cloud Generative AI Leader
- Machine Learning Specialization — Stanford University & DeepLearning.AI