Sai Narayana B

+91 8248881924 sainarayana2311@gmail.com GitHub: SaiNarayana-B LinkedIn: sai23
Portfolio: sai-narayana-portfolio.vercel.app

SUMMARY

Applied AI Engineer with close to 2 years at Quantiphi taking LLM systems from prototype to production on Google Cloud. Works across the full loop: agents and tool integrations, the evaluation and usage analytics that show whether they work, and the cost and latency engineering that keeps them viable at scale. As a Forward Deployed Engineer, embedded with a client team and used self-built spec-driven development harnesses to deliver a project scoped for 15 people with 6, in 50 of 90 planned days. Claude Certified Developer (985/1000) and 3x Google Cloud certified.

EXPERIENCE

Quantiphi
Jul 2025 - Present
Machine Learning Engineer (Applied GenAI) — *Bengaluru, India*

Spec-Driven Development Harnesses & Forward Deployed Client Engagement

- Built **spec-driven development (SDD) harnesses**: agent workflows that take requirements to specs, code, tests, and deployment. Demoed end to end to a US banking client's stakeholders.
- Embedded as one of two **Forward Deployed Engineers** with a global academic publisher's developers to build a **research paper recommendation** system using the harnesses.
- Delivered with **6 people** (1 junior and 2 senior developers, 1 architect, 2 FDEs) against a planned **15-person team**, in **50 days against a 90-day plan**: a **60% smaller team** finishing **44% sooner**.

Cypher — AI Usage Analytics (Client Scripts & LLM Gateway Endpoint)

- Built Cypher scripts that run on developers' machines to capture AI interaction logs from **Copilot, Kiro, and Claude Code** across **2,000+ employees**, scoring prompt quality and effort saved per developer.
- Built a new Cypher endpoint on the company's **LLM gateway** that analyses every request from the **Codeaira web app** and **Claude Code**, categorises and rates each interaction, and classifies work as SDD or non-SDD.
- Ran it on **flex inference** to cut cost for latency-tolerant analysis, and used a **shadow deployment** to compare two models on live traffic.

TEX Agent — Multi-Agent Staffing System (RAG & Parallel Orchestration)

- Built a **multi-agent system** on **Google ADK** that spawns parallel agents per role and returns ranked candidates with reasoning, using **RAG** over **2,000+ documents** in **Vertex AI Vector Search**.
- Cut runtime by **95%** (45 min to 5 min) by replacing sequential processing with **parallel orchestration**.

NL2SQL Agent — QAssure Platform (Text-to-SQL & MCP)

- Built a conversational agent used by **50+ PMs and business leads** that turns plain-English questions into **BigQuery SQL** through an **MCP server** and returns summarised project health insights.

LLM Cost Optimisation & Codeaira Coding Assistant

- Cut daily inference cost from **$4,000 to $2,500** (37%) with **prompt caching** and **model routing** across Gemini and Claude.
- Benchmarked the internal coding assistant against **5 open-source agents**, then rewrote its **system prompts** and added slash-command skills.

Quantiphi
Jan 2025 - Jun 2025
Machine Learning Engineer Intern — *Bengaluru, India*

AI Canvas Agent & Rapid Pitch Agent (Agentic AI & Google ADK)

- Built the **AI Canvas Agent** (solution docs and decks for **4 GTM teams**) and the **Rapid Pitch Agent** (company research to pitch deck for pre-sales).
- Set up **Google ADK** and **MCP servers** connecting enterprise data to agents across **3 internal pilot projects**.

SKILLS

Languages:	Python, SQL
AI Systems:	Agentic AI, Multi-Agent Orchestration, Spec-Driven Development, RAG, NL2SQL, MCP Servers, Tool Calling, Prompt Engineering, Prompt Caching, Model Routing, LLM Evaluation, Shadow Deployment, Flex Inference
Frameworks:	Google ADK, LangChain, LangGraph, Claude Agent SDK, Anthropic API, Gemini API, OpenAI API, FastAPI
Cloud & Tools:	GCP (Vertex AI, BigQuery, GCS, Cloud Run), Vertex AI Vector Search, Git, Claude Code, GitHub Copilot, Cursor

OPEN SOURCE PROJECT

patchwork-conventions
GitHub · PyPI · MCP Registry · Jun 2026

AST-based codebase scanner that auto-generates CONVENTIONS.md for AI coding agents

- Developed **tree-sitter AST** analysis engine detecting naming conventions, import patterns, error handling, testing frameworks, and API shapes across **5 languages** with confidence scores and real code examples.
- Shipped as CLI tool, **MCP server** (8 tools), and Python API; published to **PyPI** and **MCP Registry**; integrates with Claude Code, Cursor, and GitHub Copilot.

EDUCATION

SRM Institute of Science and Technology
2021 - 2025
B.Tech in Computer Science Engineering — *CGPA: 9.17*

CERTIFICATIONS

- Claude Certified Developer (Foundations), Anthropic: Score 985/1000
- Google Cloud Forward Deployed Engineer Cohort: Gemini Enterprise Agent Platform (4-day intensive)
- Google Cloud: Professional Machine Learning Engineer · Associate Cloud Engineer · Generative AI Leader
- Machine Learning Specialization — Stanford University & DeepLearning.AI